

```
In [1]: #heart
import os
os.chdir('desktop')
```

```
In [2]: import pandas as pd
import numpy as np
```

```
In [5]: df=pd.read_csv("heart.csv")
```

```
In [6]: df.head()
```

```
Out[6]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0

```
In [ ]: /* age (Age in years)

sex : (1 = male, 0 = female)

cp (Chest Pain Type): [ 0: asymptomatic, 1: atypical angina, 2:non-anginal pain, 3: typical angina]

trestbps (Resting Blood Pressure in mm/hg )

chol (Serum Cholesterol in mg/dl)

fbs (Fasting Blood Sugar > 120 mg/dl): [0 = no, 1 = yes]

restecg (Resting ECG): [0: showing probable or definite left ventricular hypertrophy by Estes' criteria, 1: normal, 2: having ST-T

thalach (maximum heart rate achieved)

exang (Exercise Induced Angina): [1 = yes, 0 = no]
```

```
oldpeak (ST depression induced by exercise relative to rest)

slope (the slope of the peak exercise ST segment): [0: downsloping; 1: flat; 2: upsloping]

ca [number of major vessels (0-3)]

thal : [1 = normal, 2 = fixed defect, 3 = reversible defect]

target: [0 = disease, 1 = no disease] */
```

```
In [8]: #data cleaning
df.isnull().sum()
```

```
Out[8]: age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

```
In [12]: df.nunique()
```

```
Out[12]: age          41
sex           2
cp            4
trestbps     49
chol        152
fbs           2
restecg       3
thalach      91
exang         2
oldpeak      40
slope         3
ca            5
thal          4
```

```
target      2
dtype: int64
```

```
In [10]: #Error correction in thal and caa
df['ca'].unique()
```

```
Out[10]: array([2, 0, 1, 3, 4], dtype=int64)
```

```
In [ ]: #we have to replace all 4 by nan so that unique will be 4 and not 5
```

```
In [11]: df['ca'].value_counts()
```

```
Out[11]: 0    578
1    226
2    134
3     69
4     18
Name: ca, dtype: int64
```

```
In [14]: df.loc[df['ca']==4, 'ca']=np.NaN
```

```
In [16]: df.nunique()
```

```
Out[16]: age      41
sex        2
cp         4
trestbps   49
chol      152
fbs        2
restecg     3
thalach    91
exang       2
oldpeak    40
slope       3
ca         4
thal       4
target     2
dtype: int64
```

```
In [20]: #data modeling
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [26]: df.isnull().sum()
```

```
Out[26]: age          0
sex          0
cp          0
trestbps     0
chol         0
fbs         0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           18
thal         0
target       0
dtype: int64
```

```
In [29]: df = df.fillna(df.median())
df.isnull().sum()
```

```
Out[29]: age          0
sex          0
cp          0
trestbps     0
chol         0
fbs         0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

```
In [31]: #splitting
x=df.drop(['target'],axis=1)
y=df['target']
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2, random_state=0)
```

```
In [32]: #training
lg=LogisticRegression(random_state=0)
lg.fit(x_train,y_train)
```

```
y_train_pred=lg.predict(x_train)
print("Accuracy_training : ", accuracy_score(y_train_pred,y_train))
```

Accuracy_training : 0.8670731707317073

C:\Users\AKANKSHA\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
In [33]: #testing
y_test_pred=lg.predict(x_test)
print("Accuracy_training : ", accuracy_score(y_test_pred,y_test))
```

Accuracy_training : 0.8731707317073171